

Chapter 1

Functions in tidyverse

1.1 The packages dplyr and tidyr - the core of tidyverse

These two packages are an important component of the tidyverse. The package dplyr was written by Hadley Wickham to try to simplify, speed-up and regularize commands and their arguments for manipulating data frames. The functions in **dplyr** operate on a data frame (their first argument) and the output is a data frame. HW refers to these elementary functions as 'Single table verbs'. This allows us to 'pipe' a continuous sequence of commands with each one taking a data frame as its first argument and passing its output (also a dataframe) to the next command in the pipe. Functions in dplyr are designed to form simple manipulations and in order to do more complicated tasks, the user constructs a sequence of function calls. This is like a 'grammar' of manipulation of data objects. The package is quite new (2014) and still under development but has been welcomed by the R community as a potentially unifying approach. There are many different ways of manipulating data frames in R, with inconsistent organisation of arguments and outputs.

The package tidyr complements dplyr with functions that, in the words of Hadley Wickham, 'provides a bunch of tools to help tidy up your messy datasets'.

Note:

Many of these functions have corresponding functions in base R

1.1.1 Data set

We will use the data set BackPain which is a quite large data set both concerning the number of observations and 24 variables which means that it is a little difficult to get an overview. In the examples we will use a random sample (more about the **sample** function in the simulation section) of the full data set. The imported data set is a tibble so all variables are not printed. However this is enabled with the function **glimpse**.

```
BackPain<-read_csv("../data/BackPain.csv") # The resulting object is a tibble
BackPain

# A tibble: 34,122 x 25
   id residence sex      age wealthQ physical country backPain30 agegr maritalS eduS
   <dbl> <chr>   <chr>   <dbl> <chr>   <chr>   <chr>   <chr>   <chr> <chr>   <chr>
1     1 Urban Female    66 Q3    high phy~ China    No      60-69 Div/Wid~ Comp~
2     2 Urban Female    63 Q4    mod phys~ China    No      60-69 Div/Wid~ Comp~
```

```

3      3 Urban      Female      76 Q5 rich~ high phy~ China No      70-79 Div/Wid~ No p~
4      4 Rural      Male        64 Q4      high phy~ Mexico No      60-69 Married~ Comp~
5      5 Rural      Female      67 Q1 poor~ low phys~ China No      60-69 Married~ No p~
6      6 Urban      Female      74 Q4      mod phys~ China No      70-79 Div/Wid~ Comp~
7      7 Rural      Female      51 Q3      low phys~ South ~ No      50-59 Married~ No p~
8      8 Urban      Female      80 Q5 rich~ low phys~ South ~ No      80+ Div/Wid~ Comp~
9      9 Urban      Male        67 Q4      high phy~ China No      60-69 Married~ No p~
10     10 Urban      Male        75 Q5 rich~ mod phys~ China No      70-79 Married~ Comp~
# ... with 34,112 more rows, and 14 more variables: workS <chr>, bmi <dbl>, bmi4 <chr>,
#   waistc <dbl>, smoke <chr>, alcohol <chr>, arthritis <chr>, angina <chr>,
#   depression <chr>, asthma <chr>, diabetes <chr>, comorb <dbl>, disability <dbl>,
#   height <dbl>

set.seed(1001)
bP<-BackPain[sample(nrow(BackPain),10000),] # let us take a subsample
tibble::glimpse(bP) # this is a way of getting an overview of all variables

Rows: 10,000
Columns: 25
$ id          <dbl> 27043, 27471, 5323, 9068, 28828, 23775, 24138, 17074, 18398, 25789, 8~
$ residence   <chr> "Urban", "Urban", "Rural", "Rural", "Urban", "Urban", "Rural", "Urban~
$ sex         <chr> "Female", "Female", "Female", "Female", "Female", "Female", "Female", "~
$ age         <dbl> 57, 64, 60, 70, 58, 60, 67, 63, 60, 76, 71, 65, 65, 51, 52, 53, 62, 7~
$ wealthQ     <chr> "Q3", "Q3", "Q2", "Q1 poorest", "Q5 richest", "Q1 poorest", "Q2", "Q5~
$ physical    <chr> "mod phys act", "mod phys act", "high phys act", "low phys act", "mod~
$ country     <chr> "South Africa", "China", "China", "Ghana", "Mexico", "Mexico", "China~
$ backPain30  <chr> "Yes", "Yes", "No", "No", "Yes", "Yes", "Yes", "No", "No", "Yes", "No~
$ agegr       <chr> "50-59", "60-69", "60-69", "70-79", "50-59", "60-69", "60-69", "60-69~
$ maritalS    <chr> "Married/Cohab", "Married/Cohab", "Married/Cohab", "Div/Wid/Sep", "Ma~
$ eduS        <chr> "No primary", "No primary", "No primary", "No primary", "Compl Uni/Co~
$ workS       <chr> "currently not working", "never worked", "currently working", "curren~
$ bmi         <dbl> 30.46875, 24.94150, 28.21869, 16.15668, 32.83567, 31.14878, 26.15515, ~
$ bmi4        <chr> "Obese", "Normal", "Pre-Obese", "Underweight", "Obese", "Obese", "Pre~
$ waistc      <dbl> 108.0, 82.0, 80.0, 86.0, 123.0, 101.5, 89.1, 68.0, 72.0, 87.4, 94.1, ~
$ smoke       <chr> "Never/Not Daily/Current", "Never/Not Daily/Current", "Never/Not Dail~
$ alcohol     <chr> "Abstainers", "Abstainers", "Abstainers", NA, "Abstainers", "Abstaine~
$ arthritis   <chr> "yes", "yes", "no", "no", "no", "yes", "yes", "no", "no", "no", "yes"~
$ angina      <chr> "yes", "no", "no", NA, "no", "no", "no", "no", "no", "no", "yes", "ye~
$ depression  <chr> "no", "no", "no", "yes", "no", "yes", "no", "no", "no", "no", "no", "~
$ asthma      <chr> "no", "no", "no", "no", "no", "no", "no", "no", "no", "no", "no", "ye~
$ diabetes    <chr> "yes", "no", "no", "no", "yes", "no", "no", "no", "no", "no", "no", "~
$ comorb      <dbl> 2, 1, 0, 1, 1, 2, 1, 0, 0, 0, 2, 2, 0, 0, 0, 0, 2, 1, 0, 0, 1, 0, 2, ~
$ disability  <dbl> 0.000000, 13.888890, 8.333333, 66.666660, 27.777780, 11.111110, 13.88~
$ height      <dbl> NA, 156, 155, NA, 162, NA, 155, 166, 172, 159, 155, 164, 170, NA, 165~

```

There is no text variables converted to factors in the imported data set. However, it is more convenient to use factors so we change all character vectors to factors by using the verb function **mutate_if**.

```

bP <- mutate_if(bP,is.character, as.factor)
head(bP)

# A tibble: 6 x 25
  id residence sex      age wealthQ physical country backPain30 agegr maritalS eduS
  <dbl> <fct>   <fct>   <dbl> <fct>   <fct>   <fct>   <fct>   <fct> <fct>   <fct>
1 27043 Urban   Female    57 Q3      mod phys~ South ~ Yes    50-59 Married~ No pr~
2 27471 Urban   Female    64 Q3      mod phys~ China  Yes    60-69 Married~ No pr~
3 5323  Rural   Female    60 Q2      high phy~ China  No    60-69 Married~ No pr~

```

```

4 9068 Rural      Female    70 Q1 poor~ low phys~ Ghana  No      70-79 Div/Wid~ No pr~
5 28828 Urban     Female    58 Q5 rich~ mod phys~ Mexico Yes     50-59 Married~ Compl~
6 23775 Urban     Female    60 Q1 poor~ high phy~ Mexico Yes     60-69 Div/Wid~ No pr~
# ... with 14 more variables: workS <fct>, bmi <dbl>, bmi4 <fct>, waistc <dbl>,
#   smoke <fct>, alcohol <fct>, arthritis <fct>, angina <fct>, depression <fct>,
#   asthma <fct>, diabetes <fct>, comorb <dbl>, disability <dbl>, height <dbl>

```

1.2 Single table verbs

We will now look at the following single table verbs:

In tidyverse	Examples of corresponding functions in basic R
<code>select()</code>	matrix form <code>df[,column selection]</code>
<code>filter()</code>	<code>subset()</code> or in matrix form <code>df[row selection,]</code>
<code>slice()</code>	matrix form <code>df[selection by row numbers,]</code>
<code>arrange()</code>	<code>order()</code>
<code>rename()</code>	<code>names(df)[column number]<-"newvar"</code>
<code>mutate()</code>	<code>df\$newvar<-expression</code>
<code>count()</code>	<code>table()</code> or <code>xtabs()</code>
<code>group_by()</code>	?
<code>summarise()</code>	?

Some of these also have so called scoped variants of these, e.g. `mutate_if`. See help on scoped.

1.2.1 `select()`

Column selection can be done using column names (always preferred) or their numerical position.

```

bPs <- select(bP,country, residence, sex, height, disability, diabetes)

Error in select(bP, country, residence, sex, height, disability, diabetes): unused arguments
(country, residence, sex, height, disability, diabetes)

head(bPs)

# A tibble: 6 x 6
# Groups:   country, sex, residence [1]
  country sex    residence wealthQ    count mean.disability
  <fct>   <fct>   <fct>    <fct>    <int>         <dbl>
1 China  Female Rural    Q1 poorest  277          15.7
2 China  Female Rural    Q2          260          12.4
3 China  Female Rural    Q3          172          11.8
4 China  Female Rural    Q4          156           9.99
5 China  Female Rural    Q5 richest   99           7.10
6 China  Female Rural    <NA>         6          24.5

```

In next example it is using a sequence of column numbers. Observe the alternative printing by using `kable`.

```
bPs <- select(bP, 1:4)
```

```
Error in select(bP, 1:4): unused argument (1:4)
```

```
kable(head(bPs))
```

country	sex	residence	wealthQ	count	mean.disability
China	Female	Rural	Q1 poorest	277	15.683915
China	Female	Rural	Q2	260	12.403846
China	Female	Rural	Q3	172	11.757106
China	Female	Rural	Q4	156	9.989316
China	Female	Rural	Q5 richest	99	7.098765
China	Female	Rural	NA	6	24.537038

You can use the ":" between named columns, for an interval. You can delete a column by selecting it with a minus sign in front of name or number:

```
bPs <- select(bP, -c(sex, age))
```

```
Error in select(bP, -c(sex, age)): unused argument (-c(sex, age))
```

```
head(bPs)
```

```
# A tibble: 6 x 6
```

```
# Groups:   country, sex, residence [1]
```

	country	sex	residence	wealthQ	count	mean.disability
	<fct>	<fct>	<fct>	<fct>	<int>	<dbl>
1	China	Female	Rural	Q1 poorest	277	15.7
2	China	Female	Rural	Q2	260	12.4
3	China	Female	Rural	Q3	172	11.8
4	China	Female	Rural	Q4	156	9.99
5	China	Female	Rural	Q5 richest	99	7.10
6	China	Female	Rural	<NA>	6	24.5

or:

```
bPs <- select(bP, -(3:4))
```

```
Error in select(bP, -(3:4)): unused argument (-(3:4))
```

```
head(bPs)
```

```
# A tibble: 6 x 6
```

```
# Groups:   country, sex, residence [1]
```

	country	sex	residence	wealthQ	count	mean.disability
	<fct>	<fct>	<fct>	<fct>	<int>	<dbl>
1	China	Female	Rural	Q1 poorest	277	15.7
2	China	Female	Rural	Q2	260	12.4
3	China	Female	Rural	Q3	172	11.8
4	China	Female	Rural	Q4	156	9.99
5	China	Female	Rural	Q5 richest	99	7.10
6	China	Female	Rural	<NA>	6	24.5

1.2.2 filter()

With `filter()` you select **rows** of a data frame. Note the double '=' signs as usual. If you use more than one selection condition there is "and" between. NA in a filter variable is dropped.

```
bPf <- filter(bP, country == 'China' ,
             residence == 'Rural', sex=='Female')
bPf

# A tibble: 970 x 25
   id residence sex      age wealthQ physical country backPain30 agegr maritalS eduS
   <dbl> <fct>   <fct>   <dbl> <fct>   <fct>   <fct>   <fct>   <fct> <fct>   <fct>
1  5323 Rural    Female    60 Q2      high ph~ China   No      60-69 Married~ No pr~
2  24138 Rural    Female    67 Q2      mod phy~ China   Yes     60-69 Married~ No pr~
3  4626 Rural    Female    55 Q5 rich~ high ph~ China   No      50-59 Married~ No pr~
4  15300 Rural    Female    53 Q5 rich~ high ph~ China   No      50-59 Married~ Compl~
5  9083 Rural    Female    61 Q4      low phy~ China   No      60-69 Married~ Compl~
6  9827 Rural    Female    58 Q5 rich~ high ph~ China   No      50-59 Married~ No pr~
7   605 Rural    Female    50 Q2      low phy~ China   No      50-59 Married~ Compl~
8  5993 Rural    Female    52 Q2      high ph~ China   No      50-59 Married~ No pr~
9  28901 Rural    Female    68 Q1 poor~ high ph~ China   Yes     60-69 Married~ No pr~
10 18931 Rural    Female    64 Q1 poor~ high ph~ China   No      60-69 Married~ Compl~
# ... with 960 more rows, and 14 more variables: workS <fct>, bmi <dbl>, bmi4 <fct>,
# waistc <dbl>, smoke <fct>, alcohol <fct>, arthritis <fct>, angina <fct>,
# depression <fct>, asthma <fct>, diabetes <fct>, comorb <dbl>, disability <dbl>,
# height <dbl>
```

We can get rid of those pesky NA's using the `complete.cases` function.

```
bPnf <- select(bP,bmi,waistc,age, height)

Error in select(bP, bmi, waistc, age, height): unused arguments (bmi, waistc, age, height)

summary(bPnf)

Error in summary(bPnf): object 'bPnf' not found

bPf <- filter(bPnf,complete.cases(bmi,waistc,age)) # data are now also complete in bmi and waistc

Error in filter(bPnf, complete.cases(bmi, waistc, age)): object 'bPnf' not found

# complete.cases operate on variables while omit.na operates on a data frame
summary(bPf)
```

id	residence	sex	age	wealthQ
Min. : 44	Rural:970	Female:970	Min. :50.00	Q1 poorest:277
1st Qu.: 8163	Urban: 0	Male : 0	1st Qu.:55.00	Q2 :260
Median :16685			Median :60.00	Q3 :172
Mean :16617			Mean :62.16	Q4 :156
3rd Qu.:24634			3rd Qu.:68.00	Q5 richest: 99
Max. :34115			Max. :90.00	NA's : 6

physical	country	backPain30	agegr	maritalS
high phys act:404	China :970	No :648	50-59:457	Div/Wid/Sep :197
low phys act :347	Ghana : 0	Yes :285	60-69:305	Married/Cohab:772
mod phys act :202	India : 0	NA's: 37	70-79:164	Never Married: 1
NA's : 17	Mexico : 0		80+ : 44	
	Russian Federation: 0			
	South Africa : 0			

eduS	workS	bmi	bmi4
Compl Primary :161	currently not working:243	Min. :13.30	Normal :542
Compl Sec/HS : 77	currently working :502	1st Qu.:21.06	Obese : 50
Compl Uni/Coll: 0	never worked :208	Median :23.62	Pre-Obese :278
No primary :732	NA's : 17	Mean :23.76	Underweight: 57

```

3rd Qu.:26.10  NA's      : 43
Max.      :42.56
NA's      :43

waistc      smoke
Min.       : 58.20   Current Daily      : 35
1st Qu.: 76.35   Never/Not Daily/Current:918
Median : 83.40   NA's              : 17
Mean      : 83.52
3rd Qu.: 90.10
Max.      :115.40
NA's      :42

alcohol      arthritis      angina      depression
Abstainers      :840      no :730      no :647      no :920
Non-heavy/Infreq heavy/Freq heavy drinkers: 66      yes :214      yes : 90      yes : 24
NA's            : 64      NA's: 26      NA's:233      NA's: 26

asthma      diabetes      comorb      disability      height
no :917      no :919      Min.      :0.0000      Min.      : 0.000      Min.      : 80.0
yes : 27      yes : 32      1st Qu.:0.0000      1st Qu.: 2.778      1st Qu.:150.0
NA's: 26      NA's: 19      Median :0.0000      Median : 8.333      Median :153.0
Mean      :0.3924      Mean      :12.371      Mean      :152.7
3rd Qu.:1.0000      3rd Qu.:16.667      3rd Qu.:158.0
Max.      :2.0000      Max.      :86.111      Max.      :172.0
NA's      :17      NA's      :49

```

1.2.3 slice()

Rows can also be selected by position using `slice()`.

```

bPs5 <- slice(bP, 1:5)
bPs5

# A tibble: 5 x 25
  id residence sex      age wealthQ physical country backPain30 agegr maritalS eduS
  <dbl> <fct>   <fct> <dbl> <fct>   <fct>   <fct>   <fct>   <fct> <fct>   <fct>
1 27043 Urban   Female  57 Q3    mod phys~ South ~ Yes    50-59 Married~ No pr~
2 27471 Urban   Female  64 Q3    mod phys~ China  Yes    60-69 Married~ No pr~
3 5323 Rural    Female  60 Q2    high phy~ Ghana  No     60-69 Married~ No pr~
4 9068 Rural    Female  70 Q1 poor~ low phys~ Ghana  No     70-79 Div/Wid~ No pr~
5 28828 Urban   Female  58 Q5 rich~ mod phys~ Mexico Yes    50-59 Married~ Compl~
# ... with 14 more variables: workS <fct>, bmi <dbl>, bmi4 <fct>, waistc <dbl>,
# smoke <fct>, alcohol <fct>, arthritis <fct>, angina <fct>, depression <fct>,
# asthma <fct>, diabetes <fct>, comorb <dbl>, disability <dbl>, height <dbl>

bPs10 <- slice(bP, seq(1, nrow(bP), by = 10)) # Select every 10th observation
head(bPs10)

# A tibble: 6 x 25
  id residence sex      age wealthQ physical country backPain30 agegr maritalS eduS
  <dbl> <fct>   <fct> <dbl> <fct>   <fct>   <fct>   <fct>   <fct> <fct>   <fct>
1 27043 Urban   Female  57 Q3    mod phys~ South Af~ Yes    50-59 Married~ No p~
2 8120 Urban   Female  71 Q2    mod phys~ China    No     70-79 Married~ Comp~
3 26658 Rural    Female  63 Q4    high phy~ Russian ~ Yes    60-69 Married~ Comp~

```

```

4 21774 Rural      Male      58 Q4      high phy~ Ghana      Yes      50-59 Married~ No p~
5 17090 Urban      Male      54 Q3      high phy~ China       No       50-59 Married~ Comp~
6 30346 Urban      Female    65 Q4      low phys~ China       Yes      60-69 Married~ No p~
# ... with 14 more variables: workS <fct>, bmi <dbl>, bmi4 <fct>, waistc <dbl>,
#   smoke <fct>, alcohol <fct>, arthritis <fct>, angina <fct>, depression <fct>,
#   asthma <fct>, diabetes <fct>, comorb <dbl>, disability <dbl>, height <dbl>

```

1.2.4 arrange()

The function `arrange()` is used to reorder rows. You provide a column name to control the ordering; if you want to resolve ties, add more column names. Default is ascending order. Because there are many variables so we limit the example to a few of them.

```

bPa <- arrange(bP, waistc)
head(select(bPa, residence, sex, waistc, age, wealthQ), 10) # default is ascending order

Error in select(bPa, residence, sex, waistc, age, wealthQ): unused arguments (residence,
sex, waistc, age, wealthQ)

```

We can sort by more variables and choose descending order is done as follows:

```

bPa <- arrange(bP, desc(waistc), age)
head(select(bPa, residence, sex, waistc, age, wealthQ), 10)

Error in select(bPa, residence, sex, waistc, age, wealthQ): unused arguments (residence,
sex, waistc, age, wealthQ)

```

1.2.5 rename()

Renaming columns is straightforward:

```

bPr <- rename(bP, wealthQuantile = wealthQ) # New name = old name
head(bPr)

# A tibble: 6 x 25
   id residence sex      age wealthQuantile physical country backPain30 agegr maritalS
   <dbl> <fct>   <fct>   <dbl> <fct>          <fct>   <fct>   <fct>   <fct> <fct>
1 27043 Urban   Female    57 Q3          mod phys~ South A~ Yes     50-59 Married~
2 27471 Urban   Female    64 Q3          mod phys~ China   Yes     60-69 Married~
3 5323 Rural   Female    60 Q2          high phy~ China   No      60-69 Married~
4 9068 Rural   Female    70 Q1 poorest  low phys~ Ghana   No      70-79 Div/Wid~
5 28828 Urban   Female    58 Q5 richest  mod phys~ Mexico  Yes     50-59 Married~
6 23775 Urban   Female    60 Q1 poorest  high phy~ Mexico  Yes     60-69 Div/Wid~
# ... with 15 more variables: eduS <fct>, workS <fct>, bmi <dbl>, bmi4 <fct>,
#   waistc <dbl>, smoke <fct>, alcohol <fct>, arthritis <fct>, angina <fct>,
#   depression <fct>, asthma <fct>, diabetes <fct>, comorb <dbl>, disability <dbl>,
#   height <dbl>

```

1.2.6 mutate()

The function `mutate()` adds new columns which are calculated from old columns:

```
bPm <- mutate(bP, heightInches = height/2.54)
head(select(bPm, residence, sex, age, waistc, height, heightInches))
```

Error in select(bPm, residence, sex, age, waistc, height, heightInches): unused arguments (residence, sex, age, waistc, height, heightInches)

It can also be used to change/add multiple columns:

```
bPm <- mutate(bP, heightM = height/100,
              wHr = waistc/height,
              sAge = age-50)
select(bPm, residence, sex, age, height, waistc, heightM, wHr, sAge)
```

Error in select(bPm, residence, sex, age, height, waistc, heightM, wHr, : unused arguments (residence, sex, age, height, waistc, heightM, wHr, sAge)

1.2.7 count()

We can summarise individual factors with counts of their levels using the function `count()`. It reminds of `ftable` but the result is a tibble/data frame.

```
count(bP, wealthQ)
```

```
# A tibble: 6 x 2
  wealthQ      n
  <fct>    <int>
1 Q1 poorest 1889
2 Q2         1983
3 Q3         1888
4 Q4         2054
5 Q5 richest 2159
6 <NA>        27
```

```
bPc <- count(bP, country, residence, wealthQ)
bPc
```

```
# A tibble: 66 x 4
  country residence wealthQ      n
  <fct>    <fct>    <fct>    <int>
1 China  Rural    Q1 poorest    514
2 China  Rural    Q2             513
3 China  Rural    Q3             355
4 China  Rural    Q4             335
5 China  Rural    Q5 richest    194
6 China  Rural    <NA>             7
7 China  Urban    Q1 poorest    254
8 China  Urban    Q2             249
9 China  Urban    Q3             384
10 China Urban    Q4             434
# ... with 56 more rows
```

The result **bPc** is a multi-way table in long format and the resulting object is a tibble. Thus all of it will not be printed out. If you want to see all of it try `as.data.frame(bPc)`

Notice the warnings are there because there are a significant number of NA in the data.

1.2.8 base::summary()

It can sometimes be useful to use the `base::summary()` function after selecting only those columns which are of interest. For factors you get one-way frequency tables.

```
summary(select(bP, residence, country, agegr, age))

Error in select(bP, residence, country, agegr, age): unused arguments (residence, country,
agegr, age)
```

1.3 Pipes

One characteristic with tidyverse is that pipes "`%>%`" are allowed. In this section we shall look at some examples how they are used. A pipe is a link leading from one tibble (or data frame) to another via a verb (function). A quick command for "`%>%`" is "ctrl shift m".

First we want to select some variables and print the first 4 rows. Compare with the earlier examples when `bP` was an argument.

```
bP %>% select(age,sex,wealthQ,physical,bmi) %>% head(4)

Error in select(., age, sex, wealthQ, physical, bmi): unused arguments (age, sex, wealthQ,
physical, bmi)
```

We can also in the same command set the result to a data frame (or tibble) using "`<-`" or "`=`" on the left side as usual.

```
bPp<-bP %>% select(age,sex,wealthQ,physical,bmi) %>%
  filter(sex=="Female",age<60) %>% arrange(bmi)

Error in select(., age, sex, wealthQ, physical, bmi): unused arguments (age, sex, wealthQ,
physical, bmi)
```

bPp

```
# A tibble: 4 x 2
  bmi2      n
  <fct>   <int>
1 (0,20]   344
2 (20,30] 1338
3 (30,Inf]  421
4 <NA>     94
```

```
bPp<-bP %>% filter(sex=="Female",age<60) %>%
  mutate(bmi2=cut(bmi,c(0,20,30,Inf))) %>% count(bmi2)
```

bPp

```
# A tibble: 4 x 2
  bmi2      n
  <fct>   <int>
1 (0,20]   344
2 (20,30] 1338
3 (30,Inf]  421
4 <NA>     94
```

```
bPp<-bP %>% filter(sex=="Female",age<60) %>%
  mutate(bmi2=cut(bmi,c(0,20,30,Inf)),agesq=age^2) %>% slice(1:5) %>%
  select(-c(5,7:11,13:24))

Error in select(., -c(5, 7:11, 13:24)): unused argument (-c(5, 7:11, 13:24))

bPp

# A tibble: 4 x 2
  bmi2      n
  <fct>  <int>
1 (0,20]    344
2 (20,30]  1338
3 (30,Inf]   421
4 <NA>      94
```

1.4 Manipulating factors with tidyverse tools

Factors (basically categorical variables) are somewhat complicated objects but they can be very useful if we can learn how handle them.

We have already seen how to do the conversion of all character variables in a dataframe to factors:

```
bP <- mutate_if(bP, is.character, as.factor)
```

We note that this should be done with care, because there will certainly be times when conversion of character variables to factors is inappropriate - for which reason, the tidyverse dataframe functions do not automatically convert character variables to factors, which is the default in older R functions.

Here we will take a closer look at how to manipulate factors. We'll begin here by selecting a factor variable from the backPain data frame - country, in which the order of the 'levels' (categories) of the factor is not significant (nominal variable). Notice that R saves storage of factors by saving them as numerics and relating the numerics to the levels.

```
levels(bP$country)

[1] "China"          "Ghana"          "India"          "Mexico"
[5] "Russian Federation" "South Africa"

str(select(bP, country) )

Error in select(bP, country): unused argument (country)

bP %>% select(country) %>% mutate(numeric.country=as.numeric(country)) %>% slice(1:5)

Error in select(., country): unused argument (country)

# print numeric values of first 5 elements
```

Often coding for levels (input codes) is abbreviated (or sometimes more lengthy than we might like for display). By default, when plotting the level names are used on the plot. If the names are not suitable, we can then define a more suitable name for the level (we 're-code' it) and then the new code name will be used in the plots and at the same time R will re-name the level in the dataframe.

There is considerable confusion about the use of 'labels' with levels and tidyverse offers tidier (and clearer) solutions!

Basically we want to be able to do four things:

1. Coerce numeric and/or character variables to factors where it is appropriate,
2. Re-order to something more sensible. The default alphabetical order of the levels is often unsatisfactory for presentation purposes, either in a table or a plot. Use `fct_relevel`.
3. Re-name unnecessarily terse, long or meaningless names. Use `fct_recode`
4. Cut a numeric variable into named groups to create a new factor variable. Use one of:
 - `cut_interval` which makes `n` groups with equal range,
 - `cut_number` which makes `n` groups each with approximately equal numbers of observations or
 - `cut_width`, which makes groups with a specified 'bin width'

You can see examples of the cut functions [here](#).

1.4.1 Coercing numeric and character variables to a factor

Here is how to coerce a simple numeric variable. (In this case `comorb` is coded as the number of comorbidities, 0, 1, 2 - where 2 may be 2 or more.)

```
bP %>% select(comorb) %>% str()
Error in select(., comorb): unused argument (comorb)
bP %>% select(comorb) %>% mutate(comorb=as.factor(comorb)) %>% str()
Error in select(., comorb): unused argument (comorb)
```

1.4.2 Changing the codes of a factor

We sometimes need to change the codes to give us meaningful labels for presentation, we use `fct_recode`. The recoding definition in the `fct_recode` function is "new name" = "old name". We start with converting `comorb` to a factor

```
bPm<-bP %>% mutate(comorb2=as.factor(comorb))%>%
  mutate(comorb3 = fct_recode(comorb2,"None" = "0","One" = "1","Two or more"= "2"))
bPm %>% select(comorb,comorb2,comorb3) %>% slice(1:5)
Error in select(., comorb, comorb2, comorb3): unused arguments (comorb, comorb2, comorb3)
```

1.4.3 Re-coding (renaming) categories (levels) in factors

If you simply want to change the coding of some of the levels, this is also easily done with `fct_recode`. Here is the call to reduce the length of the Russian Federation code and South Africa.

```
bP %>% count(country)
```

```
# A tibble: 6 x 2
  country      n
  <fct>      <int>
1 China      3810
2 Ghana      1300
3 India      1945
4 Mexico       661
5 Russian Federation 1169
6 South Africa 1115

bPr <- bP %>%
  mutate(country = fct_recode(country,
                              "Russian Fed" = "Russian Federation",
                              "Sth Africa" = "South Africa") )

bPr %>% count(country)

# A tibble: 6 x 2
  country      n
  <fct>      <int>
1 China      3810
2 Ghana      1300
3 India      1945
4 Mexico       661
5 Russian Fed 1169
6 Sth Africa 1115
```

If you look back at the preceding script you will see that there we had the level name "Russian Federation". We've changed the associated level name to "Russian Fed" and a similar change was made to "South Africa"

1.4.4 Re-ordering factor levels

To re-order, you must specify the levels in your desired order as the arguments of `fct_relevel` following the variable name. The spelling must be exactly as in the original!

```
levels(bP$bmi4)

[1] "Normal"      "Obese"      "Pre-Obese"  "Underweight"

bPm <- bP %>% mutate(bmi4 = fct_relevel(bmi4, "Underweight", "Normal", "Pre-Obese", "Obese"))
levels(bPm$bmi4)

[1] "Underweight" "Normal"      "Pre-Obese"  "Obese"
```

1.4.5 Converting numeric to grouped factors using the cut_ functions

The function `cut()` provides a quick way of converting numeric data to grouped factors. The `forcats` package (included in `tidyverse`) simplifies and extends these calls a little.

Here's how to create 6 levels with approximately the same number of observations in each group:

```
bP %>% mutate(height6 = cut_number(height, n = 6)) %>% count(height6)

# A tibble: 7 x 2
```

```

height6      n
<fct>      <int>
1 [65,150]   1174
2 (150,156]   967
3 (156,160]  1124
4 (160,165]  1040
5 (165,170]   917
6 (170,210]   662
7 <NA>      4116

```

We use `cut_interval` to cut with approximately equal ranges:

```

bPc4 <- bP %>% mutate(height4 = cut_interval(height, n=4))
bPc4 %>% count(height4)

# A tibble: 5 x 2
  height4      n
  <fct>    <int>
1 [65,101]     65
2 (101,138]    111
3 (138,174]  5264
4 (174,210]   444
5 <NA>      4116

```

... and `cut_width` let's you choose the range.

```

bPc <- bP %>% mutate(height4 = cut_width(height, width = 20))
summary(bPc$height4)

 [50,70]  (70,90]  (90,110] (110,130] (130,150] (150,170] (170,190] (190,210]  NA's
      1         43         28         64        1038         4048         635         27    4116

```

If you want to choose your own breaks, you must revert to the base R function `cut`. We can also add labels to all the `cut_` functions above.

```

bPc <- bP %>% mutate(height4 = cut(height, breaks = c(0,120,150,170,Inf)))
summary(bPc$height4)

 (0,120] (120,150] (150,170] (170,Inf]  NA's
      110      1064      4048      662    4116

bPc <- bP %>% mutate(height4 = cut(height, breaks = c(0,120,150,170,Inf),
                                labels = c("Very Short", "Short", "Average", "Tall" )))
bPc %>% count(height4)

# A tibble: 5 x 2
  height4      n
  <fct>    <int>
1 Very Short  110
2 Short     1064
3 Average   4048
4 Tall      662
5 <NA>     4116

bPc %>% mutate(height4=fct_explicit_na(height4, na_level = "(Missing)")) %>% count(height4)

# A tibble: 5 x 2

```

```

height4      n
<fct>      <int>
1 Very Short  110
2 Short      1064
3 Average    4048
4 Tall       662
5 (Missing)  4116

```

..and finally, here's how to combine levels. Simply recode grouped levels to the same name:

```

bPc6<-bP %>% mutate(height6 = cut_number(height, n = 6,
labels = c("Very Short", "Short", "Average", "Tall", "Very Tall", "Extremely Tall")))

summary(bPc6$height6)

##      Very Short      Short      Average      Tall      Very Tall Extremely Tall
##      1174      967      1124      1040      917      662
##      NA's
##      4116

bP64 <- bPc6 %>% mutate(h6_to_4 = fct_recode(height6,
                                             "short" = "Very Short",
                                             "short" = "Short",
                                             "very tall" = "Very Tall",
                                             "very tall" = "Extremely Tall"))

xtabs(~h6_to_4+height6,data=bP64)

##      height6
## h6_to_4  Very Short Short Average Tall Very Tall Extremely Tall
## short      1174  967      0      0      0      0
## Average      0      0  1124      0      0      0
## Tall      0      0      0  1040      0      0
## very tall      0      0      0      0  917  662

```

The package `forcats` is also very useful for dealing with factors. See Chapter 15 of 'R for data science' [Wickham2017] for detailed examples and extensions on the use of the functions we have seen here.

1.5 group_by()

The `summarise()` function and those we have discussed above, become much more powerful when we use grouping operations with the verb/function `group_by()`.

The other verbs are affected by grouping as follows:

- Grouped `select()` is the same as ungrouped `select()`, excepted that grouping variables are always retained.
- Grouped `arrange()` orders first by grouping variables
- The `slice()` function extracts rows within each group.
- The `count()` function counts the number of rows with each unique value of variable, so it is particularly useful for counting the frequency of levels in factors.
- The `summarise()` function is particularly useful when applied to grouped variables, and is explained in detail below.

1.5.1 Summarising groups

In summarising groups we can add columns containing the statistics (mean, sd, max, IQR etc) for every group combination of the set specified. In our bP data, we group using factor variables.

```
bPs <- bP %>%
  group_by(residence) %>%
  summarise(mean.disability = mean(disability))
bPs

# A tibble: 2 x 2
  residence mean.disability
  <fct>         <dbl>
1 Rural         19.4
2 Urban         16.0
```

Here's a further example in which we add more groups and statistics for more variables.

```
bPs <- bP %>% group_by(country, residence, sex) %>%
  summarise(mean.disability = mean(disability), disIQR = IQR(disability),
    Bmi = mean(bmi, na.rm=T))

'summarise()' has grouped output by 'country', 'residence'. You can override using the
'.groups' argument.

bPs

# A tibble: 24 x 6
# Groups:   country, residence [12]
  country residence sex    mean.disability disIQR    Bmi
  <fct>    <fct>   <fct>         <dbl>   <dbl> <dbl>
1 China   Rural   Female         12.4    13.9  23.8
2 China   Rural   Male           8.57    11.1  23.1
3 China   Urban   Female          8.16    11.1  24.8
4 China   Urban   Male           6.62     8.33  24.6
5 Ghana   Rural   Female         27.4    25    22.2
6 Ghana   Rural   Male          19.2    25.0  21.3
7 Ghana   Urban   Female         25.9    30.6  25.3
8 Ghana   Urban   Male          18.9    30.6  24.4
9 India   Rural   Female         33.6    25    20.7
10 India  Rural   Male          24.9    25    19.8
# ... with 14 more rows
```

Observe the na.rm=T which was necessary because of missing values.

We may want to change the order in the table.

```
bPs <- bP %>% group_by(sex, country, residence) %>%
  summarise(mean.disability = mean(disability), disIQR = IQR(disability), Bmi = mean(bmi, na.rm=T))

'summarise()' has grouped output by 'sex', 'country'. You can override using the '.groups'
argument.

bPs

# A tibble: 24 x 6
# Groups:   sex, country [12]
  sex    country    residence mean.disability disIQR    Bmi
```

```

  <fct> <fct>          <fct>          <dbl> <dbl> <dbl>
1 Female China         Rural          12.4  13.9  23.8
2 Female China         Urban           8.16  11.1  24.8
3 Female Ghana         Rural          27.4   25   22.2
4 Female Ghana         Urban          25.9  30.6  25.3
5 Female India         Rural          33.6   25   20.7
6 Female India         Urban          27.8  27.8  22.4
7 Female Mexico        Rural          22.2  27.8  27.6
8 Female Mexico        Urban          19.7  25.0  29.1
9 Female Russian Federation Rural      22.1  25.0  30.9
10 Female Russian Federation Urban      23.8  27.8  29.3
# ... with 14 more rows

```

..or filter to look at only one country.

```

bPs <- bP %>%
  filter(country== "China")%>%
  group_by(sex,residence) %>%
  summarise(mean.disability = mean(disability), disIQR = IQR(disability),
    Bmi = mean (bmi,na.rm=T))

```

'summarise()' has grouped output by 'sex'. You can override using the '.groups' argument.

```

bPs

# A tibble: 4 x 5
# Groups:   sex [2]
  sex    residence mean.disability disIQR    Bmi
<fct> <fct>          <dbl> <dbl> <dbl>
1 Female Rural          12.4  13.9  23.8
2 Female Urban           8.16  11.1  24.8
3 Male   Rural           8.57  11.1  23.1
4 Male   Urban           6.62   8.33  24.6

```

When groups vary significantly in size it is prudent to include counts of observations.

```

bPs <- group_by(bP, country, sex, residence, wealthQ) %>%
  summarise(count=n(), mean.disability=mean(disability))

```

'summarise()' has grouped output by 'country', 'sex', 'residence'. You can override using the '.groups' argument.

```

bPs

# A tibble: 130 x 6
# Groups:   country, sex, residence [24]
  country sex    residence wealthQ    count mean.disability
<fct> <fct> <fct>    <fct>    <int>    <dbl>
1 China  Female Rural    Q1 poorest  277      15.7
2 China  Female Rural    Q2         260      12.4
3 China  Female Rural    Q3         172      11.8
4 China  Female Rural    Q4         156       9.99
5 China  Female Rural    Q5 richest   99       7.10
6 China  Female Rural    <NA>         6      24.5
7 China  Female Urban    Q1 poorest  137      16.7
8 China  Female Urban    Q2         146      10.8
9 China  Female Urban    Q3         210       8.90
10 China Female Urban    Q4         237       5.92
# ... with 120 more rows

```


1.6 left_join(): Merging two data frames

There is a set of functions in dplyr for merging data frames. Here we'll just demonstrate the left_join() function.

We'll first create a data set iDf with individual ID (id) and a household ID (hhID). Then we create a second data set on households (hhDf) which will relate to the the first data set through the household ID (hhID) variable.

```
ID <- 1:15
hhID <- c(1,1,1,1,1,2,2,3,3,4,4,4,4,4)
iData1 <- LETTERS[1:15]
iData2 <- letters[12:26]

iDf <- data.frame(id = as.factor(ID),
                  hhID = as.factor(hhID),
                  iD1 = iData1,
                  iD2 = iData2)

iDf

##      id hhID iD1 iD2
## 1    1    1  A   l
## 2    2    1  B   m
## 3    3    1  C   n
## 4    4    1  D   o
## 5    5    1  E   p
## 6    6    2  F   q
## 7    7    2  G   r
## 8    8    3  H   s
## 9    9    3  I   t
## 10  10    3  J   u
## 11  11    4  K   v
## 12  12    4  L   w
## 13  13    4  M   x
## 14  14    4  N   y
## 15  15    4  O   z

hhID <- 1:4
hData1 <- c("X1", "X2", "X3", "X4")
hData2 <- letters[5:8]
hhDf <- data.frame(hhID = as.factor(hhID),
                  hD1 = hData1,
                  hD2 = hData2)

hhDf

##      hhID hD1 hD2
## 1      1  X1  e
## 2      2  X2  f
## 3      3  X3  g
## 4      4  X4  h
```

Now let's see if we can create a combined data frame in which the individual data frame rows are maintained, but have added to them the variables from the household data frame with values of those variables corresponding to the household listed in the individual's data frame.

We'll firstly try a left_join using the common household ID (hhID) as the 'key'. Notice how data in each household are repeated for individuals in the same household.

```
merged <- left_join(iDf, hhDf, by="hhID")
merged
```

##	id	hhID	iD1	iD2	hD1	hD2
## 1	1	1	A	l	X1	e
## 2	2	1	B	m	X1	e
## 3	3	1	C	n	X1	e
## 4	4	1	D	o	X1	e
## 5	5	1	E	p	X1	e
## 6	6	2	F	q	X2	f
## 7	7	2	G	r	X2	f
## 8	8	3	H	s	X3	g
## 9	9	3	I	t	X3	g
## 10	10	3	J	u	X3	g
## 11	11	4	K	v	X4	h
## 12	12	4	L	w	X4	h
## 13	13	4	M	x	X4	h
## 14	14	4	N	y	X4	h
## 15	15	4	O	z	X4	h

There is also a base R function **base::merge** that can be used to merge (link) data in a similar way.

1.7 Tidyr functions

1.7.1 Functions for converting between long and wide format

Since data can be stored in different ways there is sometimes a need to convert the data to the desired form. Two types of storage is wide and long format. Let us think of an example where timber volume (cubic meters) is measured with three different methods in three areas. The data can then be stored as wide format where the volume for each method is represented as separate variables (vectors). The functions we will use are from the tidyr package.

```
[1] "Wide format"
  area method1 method2 method3
1    1      210     242     207
2    2      135     135     111
3    3      187     201     214
```

In long format the vectors for each of the three methods are stacked. This means that there will be three times as many rows in the new data frame. We now only need one column for the volume values but we also need a column (here called method) with information on the type of method.

```
[1] "Long format"
  area method volume
1    1      1     210
2    2      1     135
3    3      1     187
4    1      2     242
5    2      2     135
6    3      2     201
7    1      3     207
8    2      3     111
9    3      3     214
```

The functions used are **gather** and **spread** which can be thought of as inverse functions of each other. Observe that we need to give two new variable names: **key** which holds the information of the kind of data (methodx) here chosen as method and **value** (volume of methodx) here chosen as volume while area is kept.

```
xwide<-data.frame(area=c(1:3),method1=c(210,135,187),method2=c(242,135,201),
                  method3=c(207,111,214))
xwide
```

	area	method1	method2	method3
1	1	210	242	207
2	2	135	135	111
3	3	187	201	214

```
xlong<-xwide %>% gather(key=method,value=volume,method1,method2,method3)
xlong
```

	area	method	volume
1	1	method1	210
2	2	method1	135
3	3	method1	187
4	1	method2	242
5	2	method2	135
6	3	method2	201
7	1	method3	207
8	2	method3	111
9	3	method3	214

```
xlong2<-xwide %>% gather(key=method,value=volume,-area)
# Alternative for the same result, all variables stacked except area
xlong %>% spread(key=method,value=volume)
```

	area	method1	method2	method3
1	1	210	242	207
2	2	135	135	111
3	3	187	201	214

...and so we are back at the data frame we started with.

1.7.2 Example - the dataset Subliminal

Let us first import the dataset. Data in the dataset Subliminal are from an intervention study where 18 students were randomized to receive either of two messages with the intention to see if this would affect their performance on the mathematics exam. The control group received neutral messages whereas the intervention group received messages confirming their learning process. All students participated in a summer school in mathematics and were tested at the beginning and end of the intervention. The dataset contains the following variables:

Message:	If the student received neutral or confirmatory messages
Before:	Test result at the beginning of the study
After:	Test result at the end of the study
Improvement:	Improvement of their results (After-Before)

Observe that we need to give two new variable names **key** which holds the information of the kind

of data (variable) here chosen as Time and **value** here chosen as Result while Message, Improvement and ind.nr is kept.

```
library(haven)
Subliminal <- read_sav("../data/Subliminal.sav")

sub<-Subliminal %>% cbind(ind.nr=1:nrow(Subliminal)) # add individual number
sub %>% slice(1:5)
```

	Message	Before	After	Improvement	ind.nr
1	positive	18	24	6	1
2	positive	18	25	7	2
3	positive	21	33	12	3
4	positive	18	29	11	4
5	positive	18	33	15	5

```
nrow(sub)

[1] 18

sub_lf<-sub %>% gather(key=Time,value=Result,Before,After)
sub_lf %>% arrange(ind.nr) %>% filter(ind.nr<=5)
```

	Message	Improvement	ind.nr	Time	Result
1	positive	6	1	Before	18
2	positive	6	1	After	24
3	positive	7	2	Before	18
4	positive	7	2	After	25
5	positive	12	3	Before	21
6	positive	12	3	After	33
7	positive	11	4	Before	18
8	positive	11	4	After	29
9	positive	15	5	Before	18
10	positive	15	5	After	33

```
nrow(sub_lf)

[1] 36

sub_wf<-sub_lf %>% spread(key=Time,value=Result) %>% arrange(ind.nr)

head(sub_wf)
```

	Message	Improvement	ind.nr	After	Before
1	positive	6	1	24	18
2	positive	7	2	25	18
3	positive	12	3	33	21
4	positive	11	4	29	18
5	positive	15	5	33	18
6	positive	16	6	36	20

```
nrow(sub_wf)

[1] 18
```

In fact the different variables that is gathered do not have to be repeated measurements of the same kind. It can be completely different measures. Let us see an example using the sample of the BackPain data.

```
bPg<- bP %>% gather(key=var,value=value,disability,bmi,age)
bPg %>% count(var)

# A tibble: 3 x 2
  var          n
  <chr>      <int>
1 age      10000
2 bmi      10000
3 disability 10000

# the first three rows in each group
bPg %>% group_by(var) %>% select(var,value,residence,sex,wealthQ,physical,country) %>% slice(1:3)

Error in select(., var, value, residence, sex, wealthQ, physical, country): unused arguments
(var, value, residence, sex, wealthQ, physical, country)
```

Own experimentation

The last tibble bP is cut to show only the first three rows per group. Check if it looks ok also further down. Can you restore it to the original by use of spread? Try out other combinations of variables when using gather.